
From H. R. 9510 to Federal Law: A Narrative Case for Verified Physical AI in the Nation’s Oncology Trials

Kevin Kawchak 
CEO ChemicalQDevice
10.5281/zenodo.xxxxxxxx

June 13, 2026 (draft-narrative, v0.1.0)

Abstract.

[DRAFTING INSTRUCTIONS: write a single paragraph under 250 words, no citations. Process the thesis from `review/prompts/prompt-narrative.md` (the eight SECTION GOALS) and the evidence base in `cancer-automated/papers/VVUQ-02` (ten-gate VVUQ suite, 172 of 172 tests). State the three guiding questions and the single ask: enact H. R. 9510 v5.0 (Kawchak2026HR9510Bill, DOI 10.5281/zenodo.20619762) into Federal law.]

Disclaimer: This work is independent and not endorsed or sponsored by trial sponsors, FDA, CROs, sites, IRBs, regulators, or medical societies; and was generated using Artificial Intelligence.

Keywords Physical AI | Oncology clinical trials | Verification before generation | VVUQ | Medical AI legislation | H. R. 9510 | Patient safety | Emotional narrative advocacy

1 Introduction: a quiet question in the committee room

This review is written for a reader who shapes medical and artificial intelligence law but who has never driven a surgical robot or written a line of verification code. It therefore defines its terms as it goes, and it asks one quiet question in eight different ways: how did we ever ask human beings, alone, to carry a task this large without help?

A few definitions make the rest of the story readable. *Physical AI* is artificial intelligence that acts in the physical world through a machine, such as a surgical humanoid in an oncology trial, rather than only returning text on a screen. *Verification before generation* is the principle at the center of H. R. 9510: a machine may propose an action, but the proposal is checked against fixed, published standards before anything touches a patient. *Claude Code* is the software agent that writes the candidate procedure; *Codex* is a second, independent agent that reviews it; and *VVUQ* (verification, validation, and uncertainty quantification) is the ten-gate examination the proposal must pass, resolving to ACCEPT, ESCALATE to a qualified human, or BLOCK before execution.

[DRAFTING INSTRUCTIONS: in the full-narrative, expand these definitions into a full page using the exact sources: (a) the bill itself, `review/references/author_works.bib`

entry Kawchak2026HR9510Bill (DOI 10.5281/zenodo.20619762); (b) the verification evidence in cancer-automated/papers/VVUQ-02 (the ten gates, 172 of 172 automated tests, the ACCEPT / ESCALATE / BLOCK decisions); (c) the colored Mermaid figures in review/mermaid/diagrams/01-verification-before-generation.md and 02-eight-emotional-pillars.md. Define each lay term the first time it appears and state the three guiding questions verbatim.]




Figure 1. Verification before generation, in brief. The full figure (Mermaid 01, reproduced as colored TikZ) adds the ACCEPT, ESCALATE, and BLOCK outcomes in the full-narrative stage.

The three guiding questions return throughout: How did we ever utilize humans alone for this challenging task? Why did we ever fully trust the prior human clinical trial system if there were so many opportunities for compounding human error? And why would the human peer-review process continue to be used for increasingly complex and voluminous AI outputs? The eight sections that follow answer them, in turn, through compassion, fear of preventable harm, moral outrage, hope, responsibility, protection, trust, and urgency.

Contents

1	Introduction: a quiet question in the committee room	1
2	Compassion and empathy: what happens if we do nothing?	3
3	Fear of preventable harm: the errors we can no longer accept	3
4	Moral outrage: this should not happen in America	4
5	Hope: the opportunity to save lives if we act	4
6	Responsibility and duty: you have the power to prevent this	5
7	Protection of vulnerable people: a duty of care, written down	5
8	Trust and reassurance: why this technology can be trusted	6
9	Urgency: patients are waiting now	7
	References	8
	Acknowledgments	11

Ethical Disclosures	11
Rights and Permissions	11
Cite This Review	12
Data Availability	12

2 Compassion and empathy: what happens if we do nothing?

Begin where every successful health law begins, with a person. A lawmaker may hear statistics from experts all day, but a patient describing a long wait for a trial slot, or a parent describing a child with a rare cancer, creates an emotional anchor that the data alone cannot. The honest question a patient story poses is the one a committee cannot avoid: what happens if we do nothing?

[DRAFTING INSTRUCTIONS: in the full-narrative, ground this section in `review/references/author_works.bib` entry `Kawchak2026CancerPatientsJourney` (A Cancer Patient’s Journey Through a Regulated Physical AI Oncology Trial, DOI 10.5281/zenodo.19119939) and pair the appeal with credible evidence from `review/references/narrative_refs.bib` entry `novak2020patient`. Reproduce Mermaid figure `review/mermaid/diagrams/11-patient-journey.md` as a colored TikZ figure. Add a body-width table of the four patient archetypes (serious illness, family caregiver, child with rare disease, person unable to access treatment) with the column format `>\raggedright\arraybackslashp{<width>}.]`




Figure 2. The compassion path (placeholder; expanded from Mermaid 11 in the full-narrative).

The point is not to substitute feeling for evidence. It is to make the evidence memorable by attaching it to the people it protects, which is exactly the balance the advocacy literature recommends.

3 Fear of preventable harm: the errors we can no longer accept

The second appeal is the fear of harm that did not have to happen. Missed diagnoses, medication errors, unsafe devices, and delayed treatment have driven drug-safety regulation, device oversight, and electronic health record requirements for a generation. The underlying message is simple: without this law, more people will suffer than need to.

[DRAFTING INSTRUCTIONS: in the full-narrative, anchor the fear appeal in the two strongest credible sources in `review/references/narrative_refs.bib`: `kohn2000toerr` (To Err Is Human) and `makary2016medical` (medical error as a leading cause of death). Contrast the legacy chain of manual handoffs with the verified path using `cancer-automated/papers/VVUQ-02` (each gate catches a defect where it occurs). Reproduce Mermaid figures `review/mermaid/diagrams/04-compounding-human-error.md` and `20-patient-throughput-flow.md` as colored TikZ. Use the KEY TERM markup to note where committee debate would harden the safety language.]




Figure 3. Compounding error versus a contained defect (placeholder; expanded from Mermaid 04 in the full-narrative).

A bill cannot promise that nothing will ever go wrong. It can promise that the system checks itself at every step, so the same mistake is not made twice and is not made invisibly.

4 Moral outrage: this should not happen in America

Many health laws gain momentum after a perceived injustice: an insurance denial, an unaffordable treatment, a discriminatory algorithm, or unequal access to care. The framing becomes a sentence a committee can repeat on the floor: this is unfair, and it should not happen here. Outrage is effective because it draws both public attention and press coverage to a specific, fixable wrong.

[DRAFTING INSTRUCTIONS: in the full-narrative, anchor the outrage appeal in review/references/narrative_refs.bib entry obermeyer2019dissecting (racial bias in a widely used clinical algorithm) so the injustice is documented, not asserted. Connect it to the bill’s transparency and bias-surveillance provisions and to review/references/author_works.bib entry Kawchak2026PatientPriorityProposedU (patient priority of proposed bills). Reproduce Mermaid figure review/mermaid/diagrams/15-vulnerable-p as colored TikZ. Add a body-width table contrasting the unverified status quo with the remedy each H. R. 9510 section provides.]



Figure 4. From an unfair gap to an audited safeguard (placeholder; expanded from Mermaid 15 in the full-narrative).

The remedy is not outrage for its own sake. It is a precise statutory check that turns a documented disparity into a measured, reported, and correctable number.

5 Hope: the opportunity to save lives if we act

Not every successful health law is built on fear. The most durable ones also offer hope: earlier cancer detection, reduced physician burnout, better rural access, and faster discovery. For medical AI, the hopeful claim is concrete. We have an opportunity to save lives if we act, and the tools to do it responsibly already exist and have been tested.

[DRAFTING INSTRUCTIONS: in the full-narrative, ground the hope appeal in review/references/narrative_refs.bib entry topol2019high (human and artificial intelligence converging in medicine) and in the author’s simulation lineage in review/references/author_works.bib(Kawchak2026AcceleratedPatientPredictionPhy

Kawchak2025AcceleratingFDAComplianceCost). Reproduce Mermaid figure review/mermaid/diagrams/06-evidence-to-law-lineage.md as colored TikZ to show that the hope rests on documented work. Add a body-width table of four near-term benefits and the bill mechanism that enables each.]




Figure 5. Hope grounded in verified tools (placeholder; expanded from Mermaid 06 in the full-narrative).

Hope, paired with evidence, gives a legislator something to vote for rather than only something to vote against.

6 Responsibility and duty: you have the power to prevent this

The fifth appeal is addressed to the legislator directly. It is not sympathy but obligation: you have the power to prevent this, and future generations will judge what we do now. This is the moment the narrative turns from describing a problem to naming the body that can solve it and the precise instrument, H. R. 9510, that lets it.

[DRAFTING INSTRUCTIONS: in the full-narrative, frame the duty appeal with review/references/narrative_refs.bib entry moreland2015hearing (how credible testimony informs the legislative record) and weave in the MAIN ACTIONS (advocacy, coalition building, policy entrepreneurship, reconciliation) and the KEY TERMS (testimony and committee hearings, markup). Reproduce Mermaid figures review/mermaid/diagrams/03-bill-to-law-process.md and 14-reconciliation-gitgraph.md as colored TikZ. Add a body-width table mapping each MAIN ACTION to the committee step where it has effect.]




Figure 6. The duty to move a bill from introduction to enactment (placeholder; expanded from Mermaid 03 in the full-narrative).

Responsibility is the bridge from feeling to action. It tells the committee that the next move is theirs and that the record will remember it.

7 Protection of vulnerable people: a duty of care, written down

The sixth appeal is protective rather than purely compassionate. Health debates return again and again to the groups most exposed to harm: children, the elderly, disabled patients, veterans, and rural communities. For medical AI the specific worries are bias against minority populations, errors that fall hardest on the vulnerable, and consent that is assumed rather than confirmed.

[DRAFTING INSTRUCTIONS: in the full-narrative, anchor the protection appeal in review/references/narrative_refs.bib entries obermeyer2019dissecting and fda2021samd (the FDA AI/ML SaMD action plan) and in the author's site-documentation and CFR adaption works (Kawchak2026PhysicalAIOncologyClinical, Kawchak2026Adaption21CFRPartb, on 21 CFR Part 50 informed consent). Reproduce Mermaid figure review/mermaid/diagrams/15-vulnerable-population-safeguards.md as colored TikZ. Add a body-width table listing each protected group, its specific risk, and the bill safeguard that addresses it.]




Figure 7. Protection written into the procedure (placeholder; expanded from Mermaid 15 in the full-narrative).

Protection that is written into statute, measured, and reported is stronger than protection that depends on the good intentions of whoever happens to be on duty.

8 Trust and reassurance: why this technology can be trusted

Medical AI bills face honest public skepticism, so the seventh appeal is reassurance. Trust is not asserted; it is built from named supports: physician endorsement, academic review, mandatory safety testing, transparency provisions, pilot results, and an audit trail that anyone can inspect. The objective is a committee that can tell constituents, with evidence, that this technology can be trusted.

[DRAFTING INSTRUCTIONS: in the full-narrative, build the trust scaffold from cancer-automated/papers/VVUQ-02 (the ten-gate suite, 172 of 172 tests, each gate bound to a published consensus standard) and review/references/narrative_refs.bib entry asme2018vv40 (ASME V&V 40-2018). Reproduce Mermaid figures review/mermaid/diagrams/10-trust-scaffolding.md, 07-ten-vvuq-gates.md, and 17-system-context.md as colored TikZ. Add a body-width table binding each gate to its standard and its pass threshold, modeled on the VVUQ table convention.]




Figure 8. Trust built from standards and an audit trail (placeholder; expanded from Mermaid 10 in the full-narrative).

Reassurance works when it is specific. A named standard and a published test result persuade where a promise of safety cannot.

9 Urgency: patients are waiting now

The final appeal is timing. Legislation accelerates when delay is shown to have a cost: patients are waiting now, lives are being lost now, and the technology is advancing faster than the rules that govern it. Urgency is only persuasive when it is paired with a credible remedy, so that a committee feels able to act rather than overwhelmed.

[DRAFTING INSTRUCTIONS: in the full-narrative, anchor urgency in `review/references/narrative_refs.b` entries `kok2013finding` (fear appeals work only when paired with self-efficacy, so always present the remedy) and `fda2021samd`. Reproduce Mermaid figures `review/mermaid/diagrams/16-urgency-capability-gap.md` and `21-national-platform-capstone.md` as colored TikZ. Close the section, and the paper, with the single call to action: enact H. R. 9510 v5.0 (DOI 10.5281/zenodo.20619762). Add a body-width table of the cost of each year of delay against the cost of acting.]




Figure 9. Closing the gap before more time is lost (placeholder; expanded from Mermaid 16 in the full-narrative).

The eight appeals converge on one request. The evidence exists, the tools are tested, and the people are waiting. The committee has the power, and now the reason, to act.

References

- [1] A. Ingber, *Regulating Emotion AI in the United States: Insights from Empirical Inquiry*. PhD thesis, Deep Blue Repositories, 2025. Working paper / Doctoral dissertation.
- [2] S. Moreland-Russell, C. Barbero, S. Andersen, N. Geary, E. A. Dodson, and R. C. Brownson, ““Hearing from all sides”: How legislative testimony influences state level policy-makers in the United States,” *International Journal of Health Policy and Management*, vol. 4, no. 2, pp. 91–98, 2015. <https://doi.org/10.15171/ijhpm.2015.13>.
- [3] L. L. Novak, S. George, K. A. Wallston, Y. A. Joosten, T. L. Israel, C. L. Simpson, and C. H. Wilkins, “Patient stories can make a difference in patient-centered research design,” *Journal of Patient Experience*, vol. 7, no. 6, pp. 1438–1444, 2020. <https://doi.org/10.1177/2374373520958340>.
- [4] M. Putturaj, ““Honouring the storyteller”: The potential of Playback Theatre in health policy and systems research,” *Health Policy and Planning*, vol. 40, no. 7, pp. 809–817, 2025. <https://doi.org/10.1093/heapol/czad000>.
- [5] G. Kok, L. K. Bartholomew, G. S. Parcel, N. H. Gottlieb, and M. E. Fernández, “Finding theory- and evidence-based alternatives to fear appeals: Intervention mapping,” *International Journal of Psychology*, vol. 49, no. 2, pp. 98–107, 2013. <https://doi.org/10.1002/ijop.12001>.
- [6] L. T. Kohn, J. M. Corrigan, and M. S. Donaldson, eds., *To Err Is Human: Building a Safer Health System*. Washington, DC: National Academies Press, 2000. Institute of Medicine. <https://doi.org/10.17226/9728>.
- [7] M. A. Makary and M. Daniel, “Medical error - the third leading cause of death in the us,” *BMJ*, vol. 353, p. i2139, 2016. <https://doi.org/10.1136/bmj.i2139>.
- [8] Z. Obermeyer, B. Powers, C. Vogeli, and S. Mullainathan, “Dissecting racial bias in an algorithm used to manage the health of populations,” *Science*, vol. 366, no. 6464, pp. 447–453, 2019. <https://doi.org/10.1126/science.aax2342>.
- [9] E. J. Topol, “High-performance medicine: the convergence of human and artificial intelligence,” *Nature Medicine*, vol. 25, no. 1, pp. 44–56, 2019. <https://doi.org/10.1038/s41591-018-0300-7>.
- [10] U.S. Food and Drug Administration, “Artificial Intelligence/Machine Learning (AI/ML)-Based Software as a Medical Device (SaMD) Action Plan.” U.S. Food and Drug Administration, 2021. <https://www.fda.gov/medical-devices/software-medical-device-samd/artificial-intelligence-and-machine-learning-software-medical-device>.
- [11] American Society of Mechanical Engineers, “Assessing Credibility of Computational Modeling Through Verification and Validation: Application to Medical Devices (ASME V&V 40-2018),” 2018. <https://www.asme.org/codes-standards/find-codes-standards/v-v-40-assessing-credibility-computational-modeling-verification-validation>.
- [12] United States Congress, “Statutory Pay-As-You-Go Act of 2010, Pub. L. 111-139,” 2010. <https://www.govinfo.gov/app/details/PLAW-111publ139>.

- [13] K. Kawchak, "H. r. 9510 (bill v5.0) 2026," jun 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [14] K. Kawchak, "H. r. 9510 (bill v4.0) 2026," jun 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [15] K. Kawchak, "H. r. 9510 (bill v3.0) 2026," jun 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [16] K. Kawchak, "Verification before generation act of 2026," jun 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [17] K. Kawchak, "Vvuq physical ai oncology trial bill," may 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [18] K. Kawchak, "Mobile pancreatic cancer unitree h2 surgical humanoid with priority vvuq," may 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [19] K. Kawchak, "Vvuq oncology clinical trial llm verification automation priority over existing generated code," may 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [20] K. Kawchak, "Threefold humanoid 24/7 adverse event oncology trial response team: 4-site rotation," may 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [21] K. Kawchak, "2030: 60 second pancreatic ductal adenocarcinoma robotic whipple procedure & daraxonrasib simulation," may 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [22] K. Kawchak, "2030: 60 second glioblastoma ai robotic surgery," may 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [23] K. Kawchak, "Patient priority of proposed u.s. bills for physical ai oncology clinical trials," may 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [24] K. Kawchak, "Accelerated patient prediction in physical ai oncology clinical trials: 4 extensive simulations," may 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [25] K. Kawchak, "Fully automated sponsor: Physical ai oncology clinical trial platform," apr 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [26] K. Kawchak, "National platform for physical ai oncology trials," mar 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [27] K. Kawchak, "Physical ai oncology clinical trial site complete documentation package," mar 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [28] K. Kawchak, "A cancer patient's journey through a regulated physical ai oncology trial," mar 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [29] K. Kawchak, "Adaption: 21 cfr part 312, end-to-end physical ai oncology clinical trial unification," mar 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [30] K. Kawchak, "Adaption: 21 cfr part 50, end-to-end physical ai oncology clinical trial unification," mar 2026. ORCID record work; source snapshot exported from ORCID webarchive.

- [31] K. Kawchak, "Adaption: Ich harmonised guideline, end-to-end physical ai oncology clinical trial unification," mar 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [32] K. Kawchak, "National mcp servers for physical ai oncology clinical trial systems," mar 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [33] K. Kawchak, "Trialmcp: Mcp servers for physical ai oncology clinical trial systems," mar 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [34] K. Kawchak, "Federated learning physical ai oncology trials unification," feb 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [35] K. Kawchak, "Unification standard level for physical ai oncology trials," feb 2026. ORCID record work; source snapshot exported from ORCID webarchive.
- [36] K. Kawchak, "Code generation competition: 16 proprietary vs. open-source llms & iterative learning based on fda adverse event reporting system," dec 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [37] K. Kawchak, "Ai peer review acceleration of llm-generated glioblastoma clinical trial patient matching ml, fda/ich/iso, and fastapi," dec 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [38] K. Kawchak, "Llm-generated glioblastoma drug synergy machine learning: From rapid code prototypes to project deliverables package," nov 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [39] K. Kawchak, "End-to-end oncology clinical trial llm efficiency for industry adoption with fda/ich regulations," oct 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [40] K. Kawchak, "Accelerating fda compliance and cost efficiency of in silico clinical trials via ai digital twin pancreatic cancer simulation," oct 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [41] K. Kawchak, "Qsp metastatic pancreatic cancer ai clinical trial simulation from protocol to prediction: Code, vvuq, and playbook," aug 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [42] K. Kawchak, "Chatgpt 100,000 patient 24-month in silico phase iii 5-arm pancreatic cancer clinical trial triplicate," jul 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [43] K. Kawchak, "End-to-end pancreatic ductal adenocarcinoma digital twin clinical trial proposals," jun 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [44] K. Kawchak, "10 year glioblastoma clinical trial meta-analyses by autonomous ai at scale. survival, hr, ae, and rob scored in ai reports and charts, including verifications," may 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [45] K. Kawchak, "Ai revolution toward the cure of lung adenocarcinoma," apr 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [46] K. Kawchak, "Autonomous llm agent and scalable reasoning llm for generating cancer drug industry cost solutions," mar 2025. ORCID record work; source snapshot exported from ORCID webarchive.

- [47] K. Kawchak, “Cost containment of global monoclonal antibody drugs and cancer clinical trials via llm focused reasoning,” feb 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [48] K. Kawchak, “Clinical decision support based on bevacizumab cancer trials and pushing the limitations of advanced llms,” jan 2025. ORCID record work; source snapshot exported from ORCID webarchive.
- [49] K. Kawchak, “Cancer vs. conversational artificial intelligence,” dec 2024. ORCID record work; source snapshot exported from ORCID webarchive.
- [50] K. Kawchak, “mab bioprocess engineering in-context table forecasts using conversational ai literature insight generations,” dec 2024. ORCID record work; source snapshot exported from ORCID webarchive.
- [51] K. Kawchak, “Monoclonal antibody bioprocess engineering advancements using conversational artificial intelligence,” oct 2024. ORCID record work; source snapshot exported from ORCID webarchive.
- [52] K. Kawchak, “Paclitaxel biosynthesis ai breakthrough,” oct 2024. ORCID record work; source snapshot exported from ORCID webarchive.
- [53] K. Kawchak, “High dimensional and complex spectrometric data analysis of an organic compound using large multimodal models and chained outputs,” sep 2024. ORCID record work; source snapshot exported from ORCID webarchive.
- [54] K. Kawchak, “Lmm spectrometric determination of an organic compound,” aug 2024. ORCID record work; source snapshot exported from ORCID webarchive.
- [55] K. Kawchak, “Lmm chemical research with document retrieval,” aug 2024. ORCID record work; source snapshot exported from ORCID webarchive.

Acknowledgments

The author acknowledges Anthropic for access to Claude Code for the instruction, drafting, and assembly of this narrative review; OpenAI for ChatGPT subject support; and Google Gemini for additional research assistance.

Ethical Disclosures

The author declares no competing interests. This is an independent narrative review and is not endorsed, sponsored, or approved by any trial sponsor, CRO, site, IRB, regulator, or medical society, nor by CFR, ICH, or FDA. H. R. 9510 is an illustrative bill number; it is not enacted law and this review is not legal advice. All supporting numbers are simulation or illustrative results unless tied to a cited source.

Rights and Permissions

This review is distributed under the Creative Commons Attribution 4.0 International License (CC BY 4.0). To view a copy, visit <https://creativecommons.org/licenses/by/4.0/>. Reproduced public-domain U.S. Government text is used under 17 U.S.C. § 105.

Cite This Review

Kawchak K. From H. R. 9510 to Federal Law: A Narrative Case for Verified Physical AI in the Nation's Oncology Trials. Zenodo. 2026; 10.5281/zenodo.xxxxxxxx.

Data Availability

All sources, figures, and prompt records supporting this review are openly available in the `kevinkawchak/law-physical-ai-trials` repository under `review/` (the prompt, the four sub-prompts, the colored Mermaid catalog, and the three narrative stages), and draw on the public records in `kevinkawchak/cancer-automated` and the author's Zenodo deposits. All data sets are synthetic or illustrative; no patient-identifiable information is included.